

DS-UA 301
Advanced Topics in Data Science
*Advanced Techniques in ML and Deep
Learning*

LECTURE 13

Parijat Dube

ONNX

- Open Neural Network Exchange
- An open format to represent traditional machine learning and deep learning models
- ONNX Features
 - Framework interoperability
 - Models trained in one DL framework to be transferred to another for inference
 - Hardware optimizations
 - ONNX-compatible runtimes and libraries designed to maximize performance on specific DL hardware
- Supported by a community of over 20 leading companies
- Visit <http://onnx.ai>

ONNX Capabilities

- Common set of operators related to ML
- Common file format for representing ML models
- Supported Tools
 - Visit <http://onnx.ai/supported-tools>
- ONNX Tutorials
 - Visit <https://github.com/onnx/tutorials>
- Model Zoo
 - A collection of pre-trained, state-of-the-art models in the [ONNX](#) format contributed by community members
 - Visit <https://github.com/onnx/models>

ONNX in Enterprise: Microsoft

ML used in many products to improve performance and productivity

Microsoft 365

 Windows

 Microsoft Dynamics 365

 Skype

 Bing

 Microsoft HoloLens

 Microsoft | Research

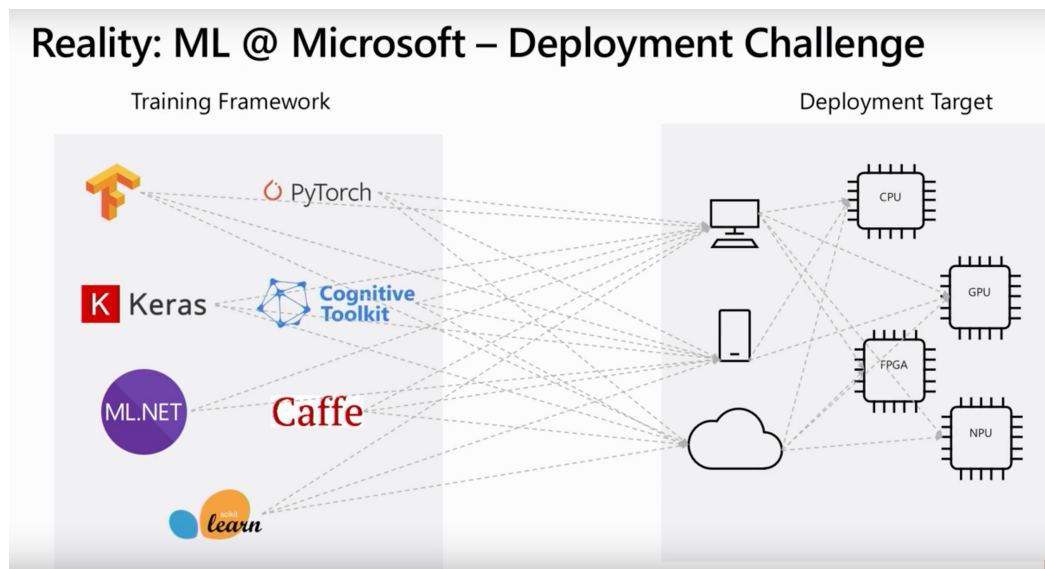
 Office 365

 XBOX

- AI developers work with different frameworks
- Deployment of trained models to production
 - Different deployment targets: cloud, IoT devices, edge devices
 - Different hardware: CPUs, GPUs, TPUs, FPGAs

Youtube Video: Open Neural Network Exchange (ONNX) in the enterprise: how Microsoft scales ML

Multiple ML frameworks + multiple deployment targets

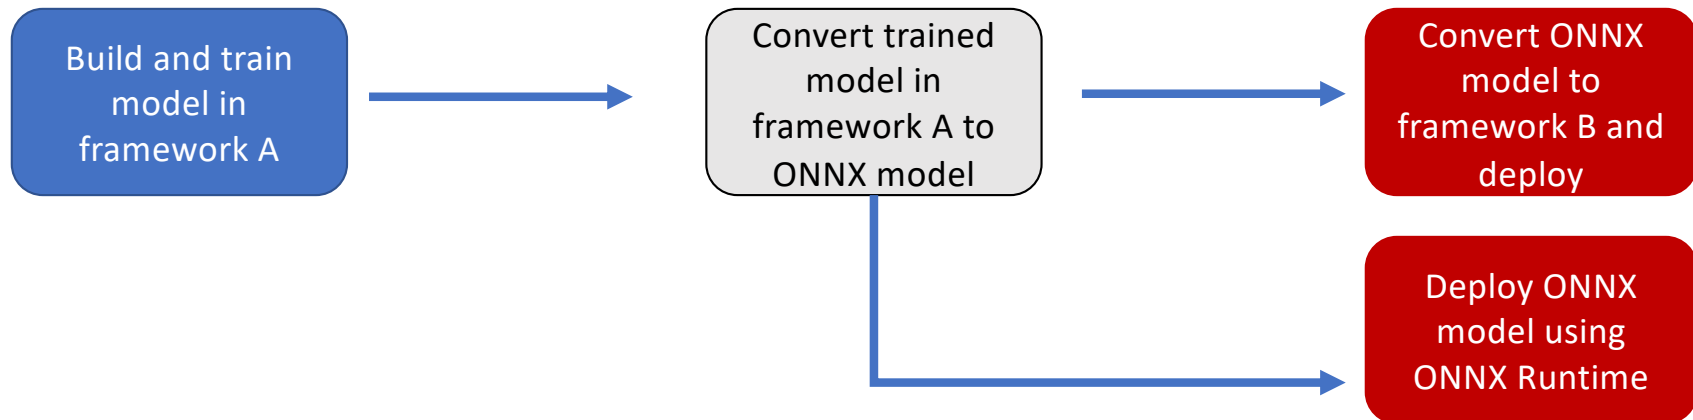


- Maintaining multiple frameworks in deployment
 - Not scalable
 - Hard to maintain
 - Degrades application performance
 - Frameworks compete for system resources
- Decouple training and deployment frameworks

Youtube Video: Open Neural Network Exchange (ONNX) in the enterprise: how Microsoft scales ML

Bridge Model Training and Deployment

Open and Interoperable industry wide standard



Example Use Case

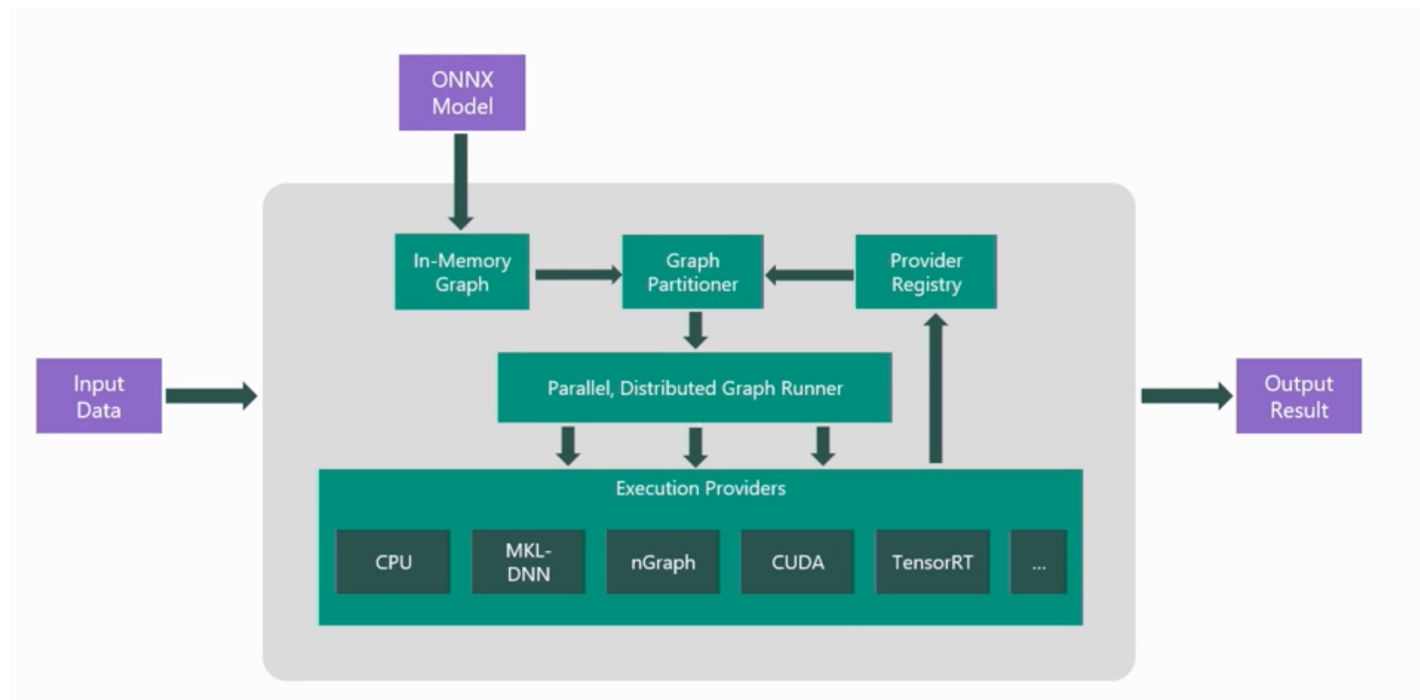
- Pytorch-ONNX-CoreML
 - <https://attardi.org/pytorch-and-coreml>

ONNX Runtime



- ONNX Runtime is a cross-platform inference and training machine-learning accelerator.
- Available on Mac, Windows, Linux platforms
- Supports full ONNX-ML spec
- Open-sourced <https://github.com/microsoft/onnxruntime>

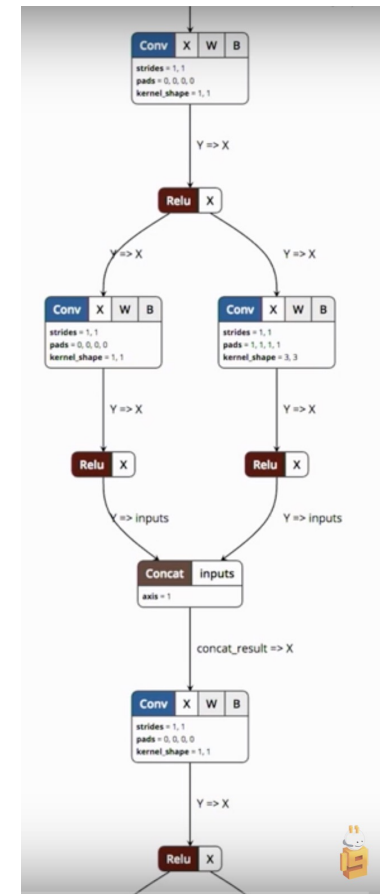
ONNX Runtime Architecture



ONNX Runtime has a graph parser which takes ONNX model, parses the graph, applies runtime optimizations like fusion ops, executes portions of the graph on specific hardware, provides the inference output result

Youtube Video: Open Neural Network Exchange (ONNX) in the enterprise: how Microsoft scales ML

Computational dataflow graph



Frameworks provide implementations of ONNX operators

Factors Contributing to Democratization of AI

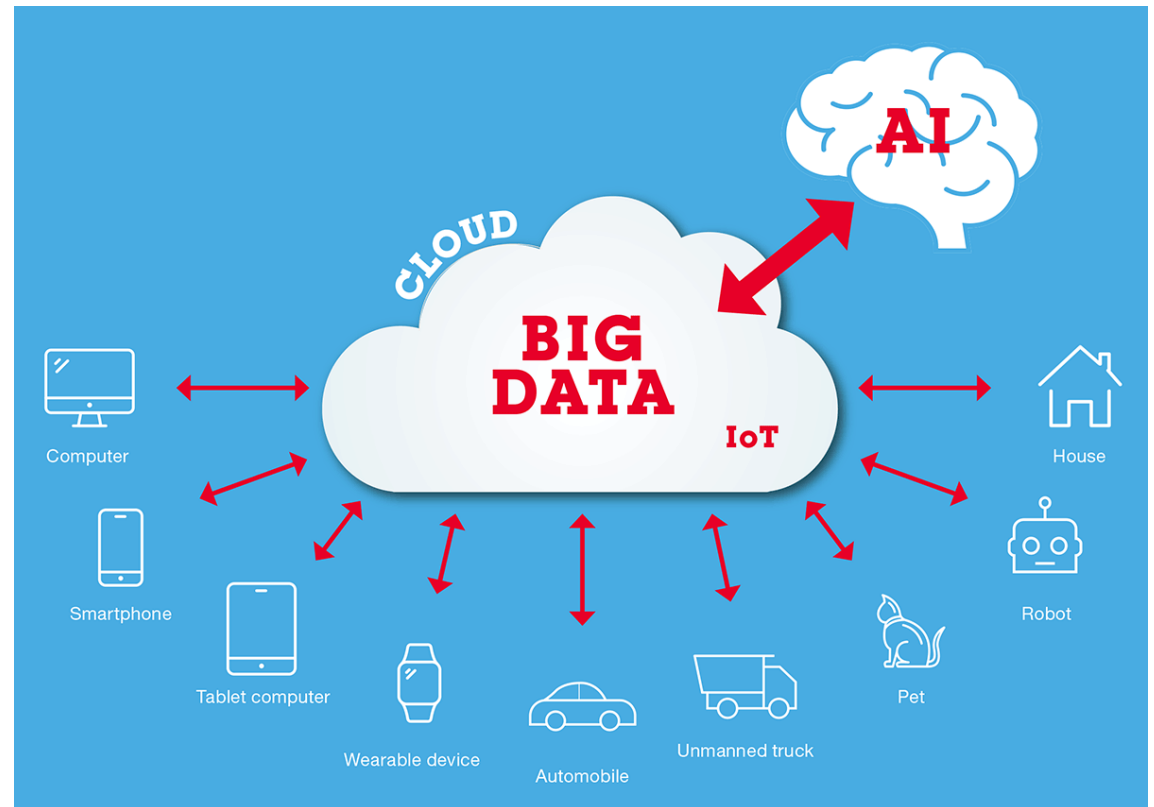
- Open benchmarking communities
 - Kaggle, Open ML
- ONNX
 - Faster and broader reuse of models; Shorten time to move to production
 - Hardware optimizations can be leveraged by many developers easily
 - Makes deep learning models portable thus preventing framework lock in
 - Greater interoperability in AI tools community
- Cloud based AI solutions

Cloud Computing

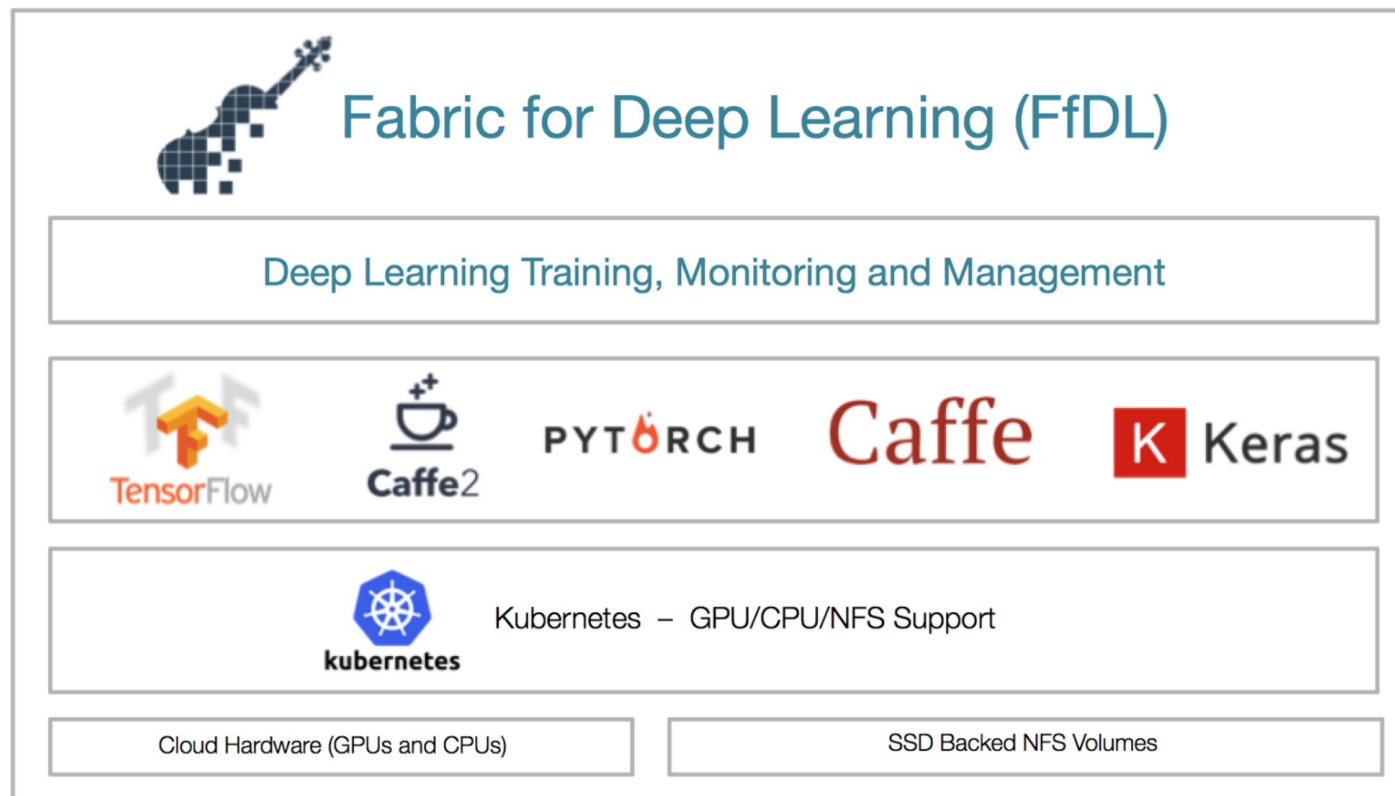
- Access to computing resources and storage on demand
- Pay-as-you go model
- Heterogeneous resources: GPUs, CPUs, storage type
- Different offering models: IaaS, PaaS, SaaS, MLaaS
- Different deployment models: Public, private, hybrid cloud
- Provisioning, maintenance, monitoring, life-cycle-management

Marriage of Cloud and AI

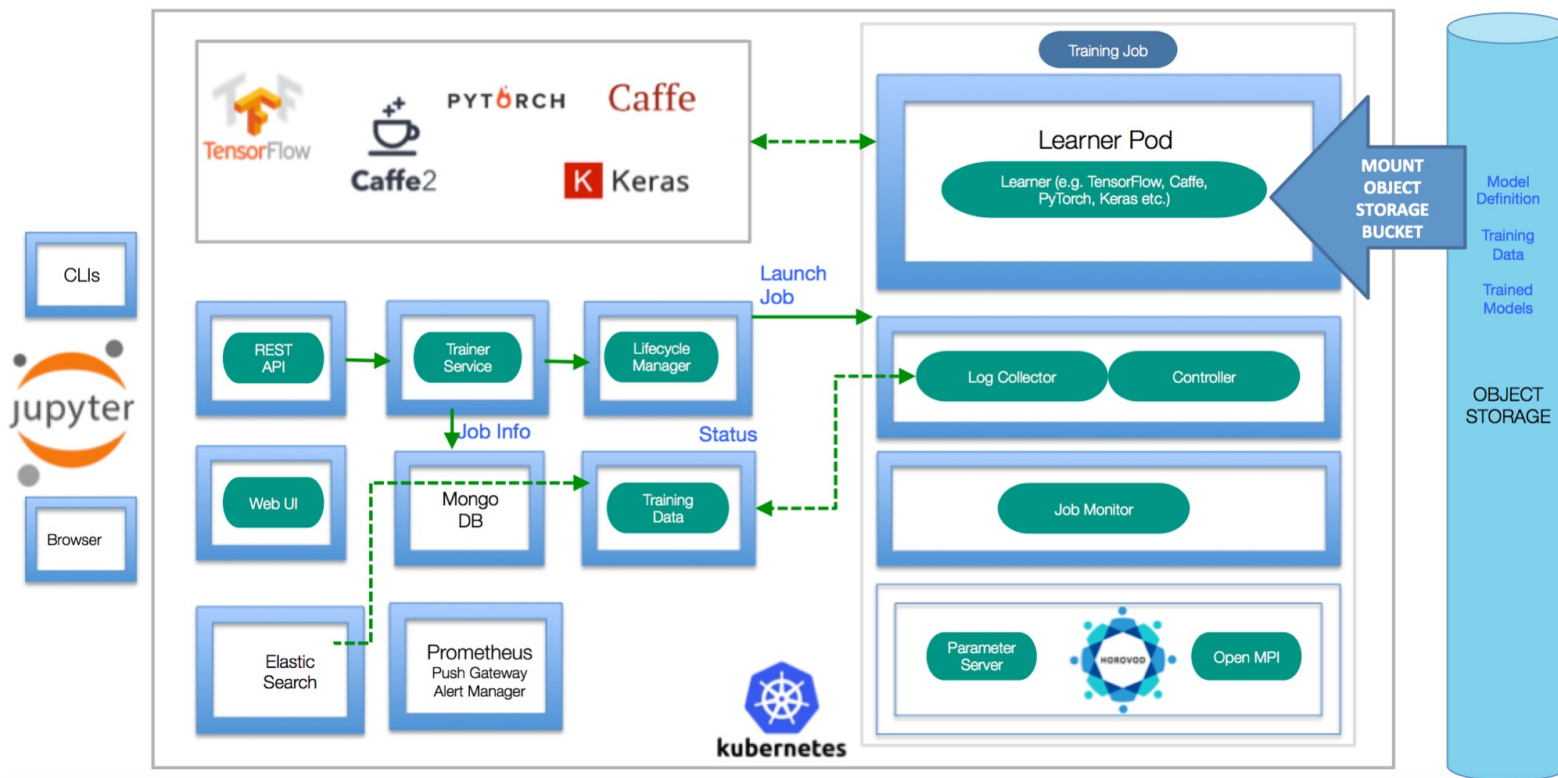
- AI
 - Harness power of Big Data and compute
- Cloud
 - Access to Big Data
 - Platform to quickly develop, deploy, and test AI solutions
 - Ease in AI reachability
- Cloud + AI is the winning combination



Deep Learning on Cloud Stack



Typical DL on Cloud architecture



Features of a Machine Learning Platform

- Frameworks: DL framework(s) supported and their version.
 - Machine learning platforms which have their own inbuilt images for different frameworks.
- Compute units: type(s) of compute units offered, i.e., GPU types.
- Model lifecycle management: tools supported to manage ML model lifecycle.
- Monitoring: availability of application logs and resource (GPU, CPU, memory) usage monitoring data to the user.
- Visualization during training: performance metrics like accuracy and throughput
- Elastic Scaling: support for elastic scaling compute resources of an ongoing job.
- Training job description: training job description file format.

Cloud based Machine Learning Services

- IBM Watson Machine Learning

<https://www.ibm.com/cloud/machine-learning>

- Amazon Sagemaker

<https://aws.amazon.com/sagemaker>

- Microsoft Azure Machine Learning

<https://azure.com/ml>

- Google Cloud Machine Learning

<https://cloud.google.com/ml-engine>

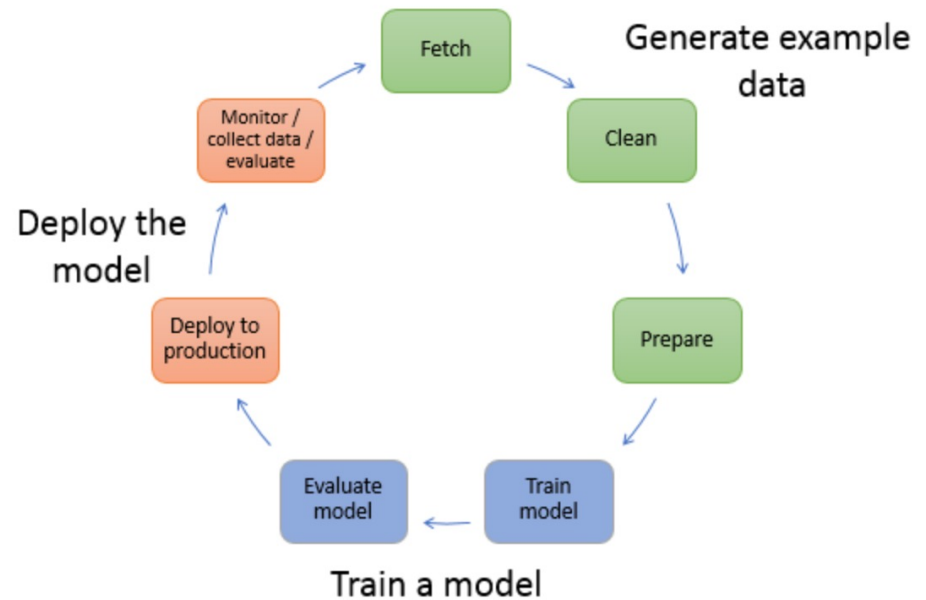
Training job specification

- YAML file
- https://dataplatform.cloud.ibm.com/docs/content/wsj/analyze-data/ml_dlaas_e2e_example.html

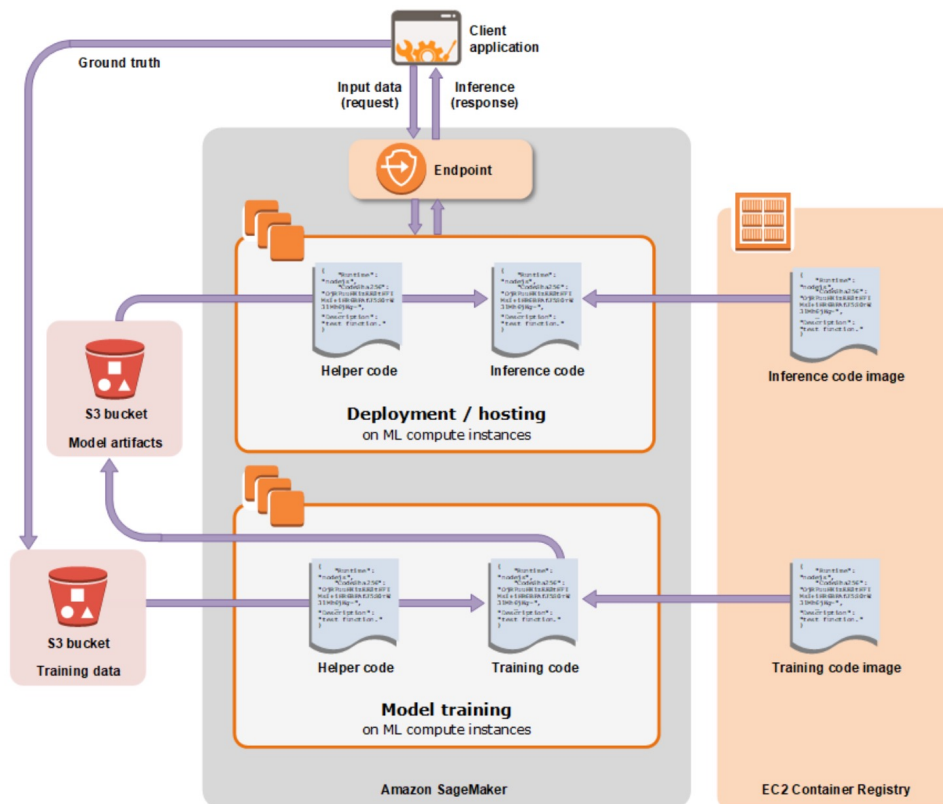
Amazon Sagemaker

- Fully managed machine learning service by Amazon
- Supports:
 - Quick and easy building and training of ML models
 - Model deployment in production-ready hosted environment

Typical ML Workflow



Train with Amazon SageMaker



Training algorithm options

- Use an out-of-the-box algorithm provided by Amazon
- Use Apache Spark MLlib with Amazon SageMaker
- Custom python code to train with DL frameworks
 - Tensorflow and Apache MXNet
- Use your own custom algorithm in any programming language and framework
- Use an algorithm from AWS Marketplace

<https://docs.aws.amazon.com/sagemaker/latest/dg/how-it-works-training.html>

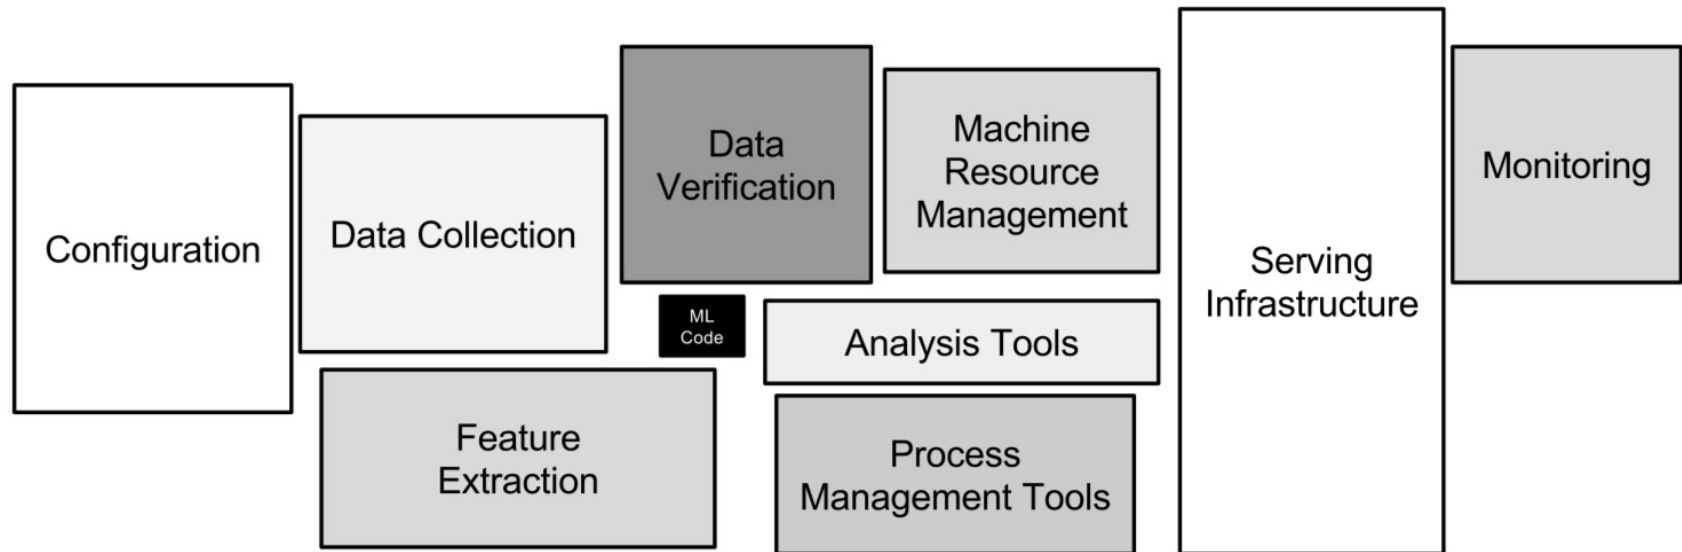
CreateTrainingJob API for Amazon Sagemaker

- https://docs.aws.amazon.com/sagemaker/latest/dg/API_CreateTrainingJob.html

Comparing Machine Learning Service Platforms

- <https://www.altexsoft.com/blog/datascience/comparing-machine-learning-as-a-service-amazon-microsoft-azure-google-cloud-ai-ibm-watson/>

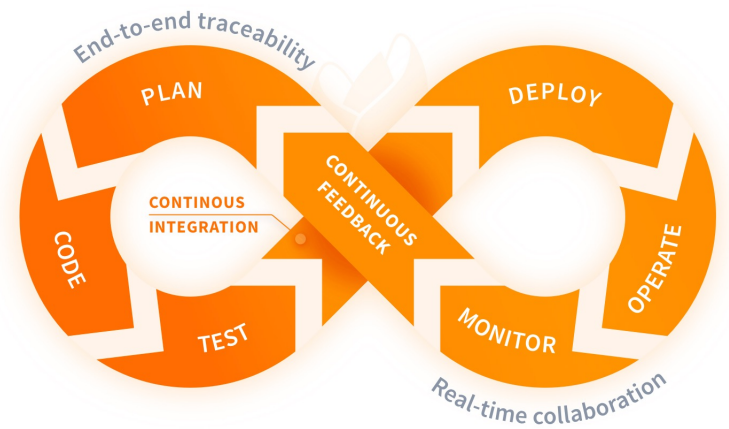
Production Grade Machine Learning Systems



The portion of ML training code in a production-grade ML system is a lot smaller than the technologies and processes needed for supporting it.

Software Engineering in ML Systems

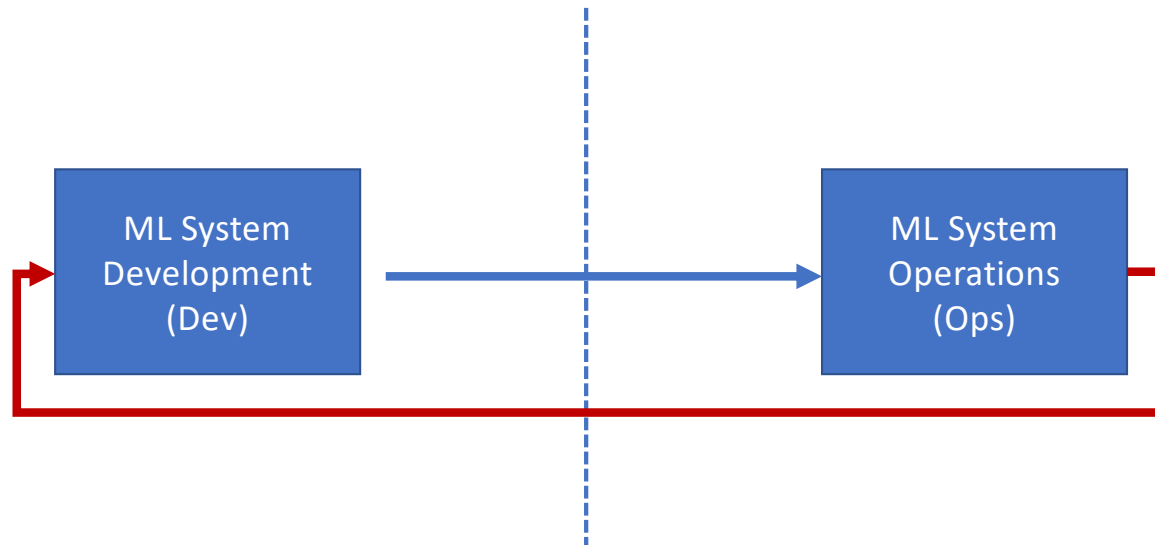
- Machine learning applications run as pipelines that ingest data, compute features, identify model(s), discover hyperparameters, train model(s), validate and deploy model(s).
- Making a model as a production-capable web service
 - Containerization (docker), cluster deployment (K8s)
 - APIs exposed as web service (Tensorflow serving/ONNX runtime)
- Workflow engines (Kubeflow) automate the ML pipeline
- Deployment monitoring and operational analytics
- Devops principles applicable to ML Systems:
 - Continuous Integration, Continuous delivery (CI/CD)
 - Predictability
 - “A model may be unexplainable—but an API cannot be unpredictable”
 - Reproducibility and Traceability
 - Provenance for Machine Learning Artifacts



ML Specific testing and monitoring apart from traditional software testing

- Data testing
- Infrastructure testing
- Model testing
- Production testing

ModelOps: The Assembly Line for ML



Goal: Accelerate model life cycle (from development to deployment)

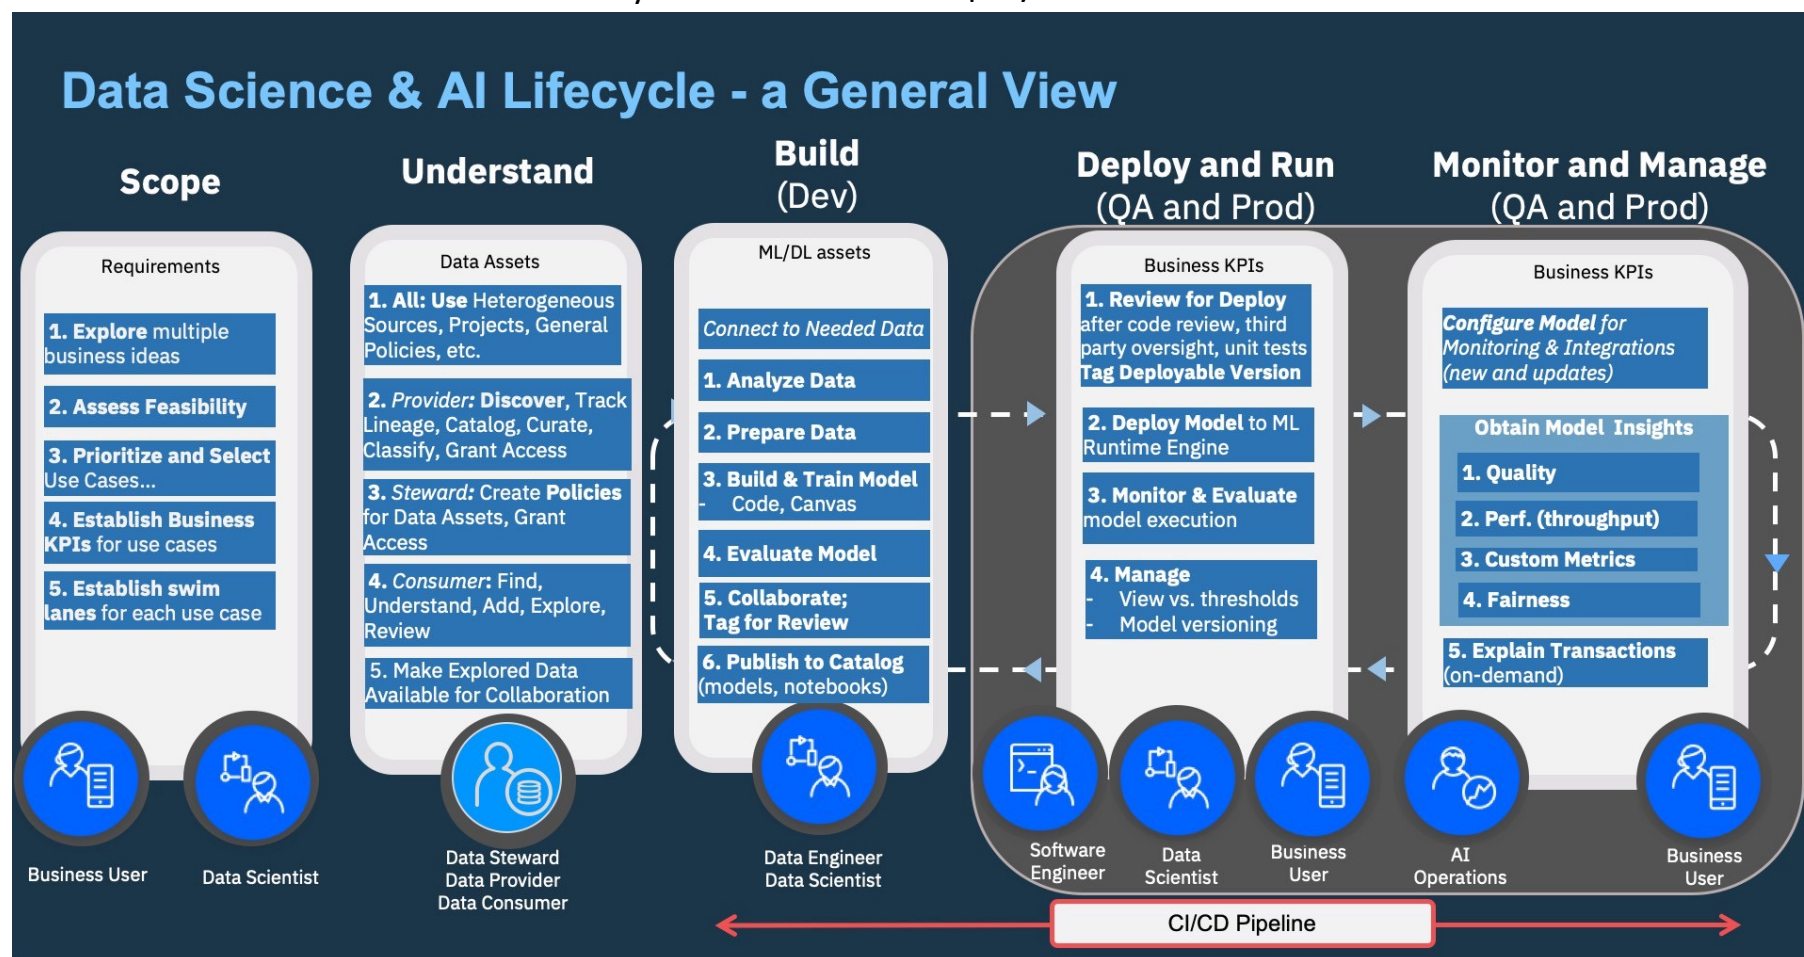
Maintain high quality model in production

Approach: automation and monitoring through development of tool-chain covering all steps of ML system construction, including development, integration, testing, releasing, deployment and infrastructure management.

Is ModelOps same as DevOps ?

Operationalizing AI

ML Models → ML Systems → Deployed ML Service → Business Value



ML specific challenges to DevOps

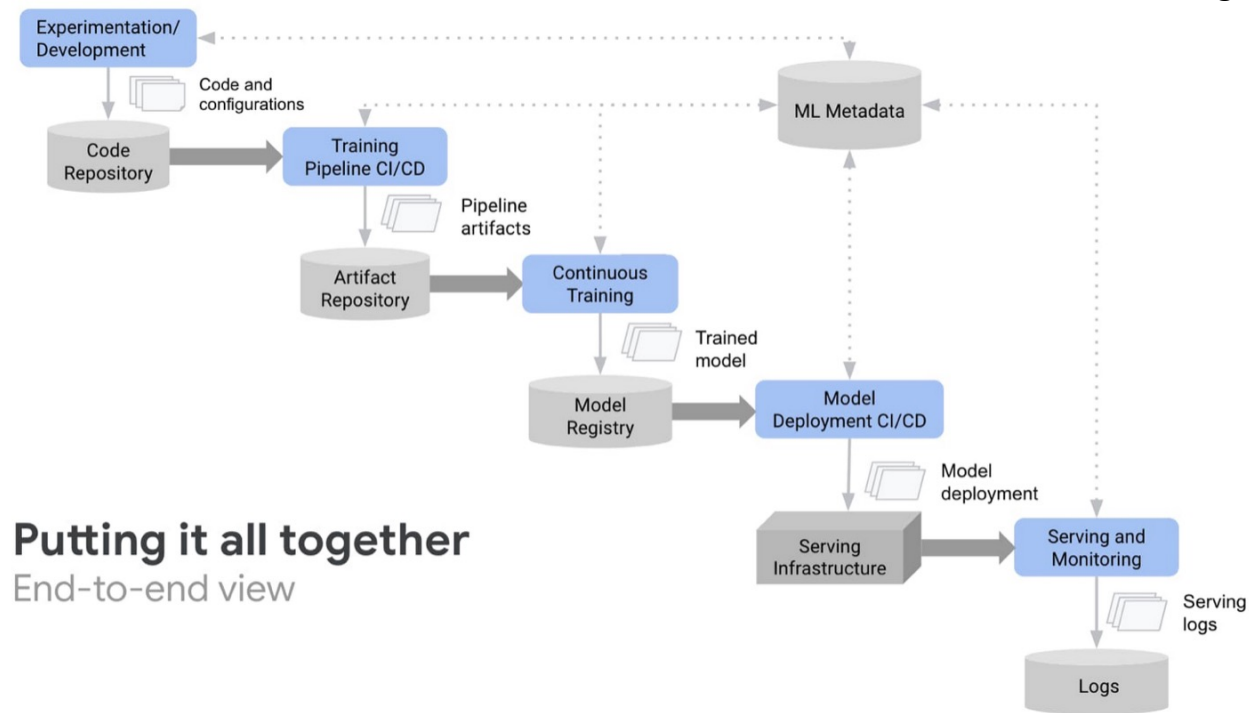
- **Continuous Integration (CI)** is not only about testing and validating code and components, but also testing and validating data, data schemas, and models.
- **Continuous Delivery (CD)** is not only about a single software package or a service, but a system (an ML training pipeline) that should automatically deploy another service (model prediction service).
- **Continuous Training (CT)** is a new property, unique to ML systems, that's concerned with automatically retraining candidate models for testing and serving.
- **Continuous Monitoring (CM)** is not only about catching errors in production systems, but also about monitoring production inference data and model performance metrics tied to business outcomes.

Google MLOps

- *MLOps* is an ML engineering culture and practice that aims at unifying ML system development (Dev) and ML system operation (Ops)
 - [An introduction to MLOps on Google Cloud](#)

MLOps with CI/CD

Repeatable and reliable pipelines
Lineage tracking of trained models



IBM Cloud Pak for Data for ML Operationalization

[Managing AI Lifecycle with MLOps](#)